

Name : Atharva Phalke

Class : D15C

Batch : C

Rollno: 65

DMBI - 3

Aim: To perform Exploratory Data Analysis and Visualization using python.

Theory:

Exploratory Data Analysis (EDA) is the process of exploring, summarizing, and visualizing data to understand its main characteristics before applying statistical models or machine learning. It helps researchers **detect underlying structures, spot anomalies, identify relationships, and test hypotheses.**

The major steps include:

1. Descriptive Statistics

- Provides numerical summaries such as mean, median, variance, min/max values, and standard deviation.
- Helps detect skewness, outliers, and unusual distributions

2. Target Variable Analysis

- The dependent variable (here: heart_disease) is visualized with count plots to check class balance.
- If the dataset is highly imbalanced, it may affect classification performance.

3. Correlation Analysis

- Pearson's correlation coefficient is calculated between numeric features.
- A heatmap helps identify strong positive/negative correlations (e.g., thalach vs. age, chol vs. bmi).
- Useful for detecting multicollinearity or redundant features.

4. Feature Distribution Analysis

- Histograms/KDE plots show how features like age, cholesterol, and bmi are distributed (normal, skewed, multimodal).
- Boxplots grouped by target show how continuous features vary between patients with and without heart disease.

5. Categorical Feature Analysis

- Countplots and bar charts show how categorical features (e.g., sex, cp, thal) are distributed across target classes.
- This helps assess the predictive power of categorical variables (e.g., chest pain type has strong association with heart disease).

6. Multivariate Visualization

- Pairplots allow simultaneous visualization of multiple variables, highlighting clusters and class separation.
- Useful to see which combinations of features separate patients with vs. without heart disease.

Importance:

- Provides deeper insight into dataset structure.
- Helps select features that are most relevant for prediction.
- Reveals outliers or errors that might need special treatment.
- Builds intuition about how independent variables influence the target outcome.

Conclusion:

- The Stroke Prediction dataset provides valuable features for building predictive models. However, proper handling of class imbalance, missing values, and outliers will be crucial for developing robust and reliable machine learning solutions.

Code and Output:

```
Exp3.py x healthcare-dataset-stroke-data.csv
Exp3.py > ...
1 # Exploratory Data Analysis (EDA) for Stroke Prediction Dataset
2
3 import pandas as pd
4 import numpy as np
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7
8 # 1. Load the dataset (update the path if needed)
9 df = pd.read_csv('C:\\Users\\athar\\OneDrive\\Desktop\\DMBI\\archive\\healthcare-dataset-stroke-data.csv') # Change filename/path as needed
10
11 # 2. Initial Data Exploration
12 print("First 5 rows:")
13 print(df.head())
14 print("\nDataset info:")
15 print(df.info())
16 print("\nStatistical summary:")
17 print(df.describe(include='all'))
18 print("\nMissing values in each column:")
19 print(df.isnull().sum())
20
```



```
First 5 rows:
   id  gender  age  hypertension  heart_disease  ...  Residence_type  avg_glucose_level  bmi  smoking_status  stroke
0  9046   Male  67.0             0              1  ...      Urban                228.69   36.6  formerly smoked      1
1  51676  Female  61.0             0              0  ...      Rural                202.21   NaN   never smoked      1
2  31112   Male  80.0             0              1  ...      Rural                105.92  32.5   never smoked      1
3  60182  Female  49.0             0              0  ...      Urban                171.23  34.4     smokes      1
4   1665  Female  79.0             1              0  ...      Rural                174.12  24.0   never smoked      1

[5 rows x 12 columns]

Dataset info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  -
0   id                    5110 non-null   int64
1   gender                5110 non-null   object
2   age                   5110 non-null   float64
3   hypertension          5110 non-null   int64
4   heart_disease         5110 non-null   int64
5   ever_married          5110 non-null   object
6   work_type             5110 non-null   object
7   Residence_type        5110 non-null   object
8   avg_glucose_level     5110 non-null   float64
9   bmi                   4909 non-null   float64
10  smoking_status        5110 non-null   object
11  stroke                5110 non-null   int64
dtypes: float64(3), int64(4), object(5)
memory usage: 479.2+ KB
None
```

```

Statistical summary:

```

	id	gender	age	hypertension	...	avg_glucose_level	bmi	smoking_status	stroke
count	5110.000000	5110	5110.000000	5110.000000	...	5110.000000	4909.000000	5110	5110.000000
unique	NaN	3	NaN	NaN	...	NaN	NaN	4	NaN
top	NaN	Female	NaN	NaN	...	NaN	NaN	never smoked	NaN
freq	NaN	2994	NaN	NaN	...	NaN	NaN	1892	NaN
mean	36517.829354	NaN	43.226614	0.097456	...	106.147677	28.893237	NaN	0.048728
std	21161.721625	NaN	22.612647	0.296607	...	45.283560	7.854067	NaN	0.215320
min	67.000000	NaN	0.080000	0.000000	...	55.120000	10.300000	NaN	0.000000
25%	17741.250000	NaN	25.000000	0.000000	...	77.245000	23.500000	NaN	0.000000
50%	36932.000000	NaN	45.000000	0.000000	...	91.885000	28.100000	NaN	0.000000
75%	54682.000000	NaN	61.000000	0.000000	...	114.090000	33.100000	NaN	0.000000
max	72940.000000	NaN	82.000000	1.000000	...	271.740000	97.600000	NaN	1.000000

```

[11 rows x 12 columns]

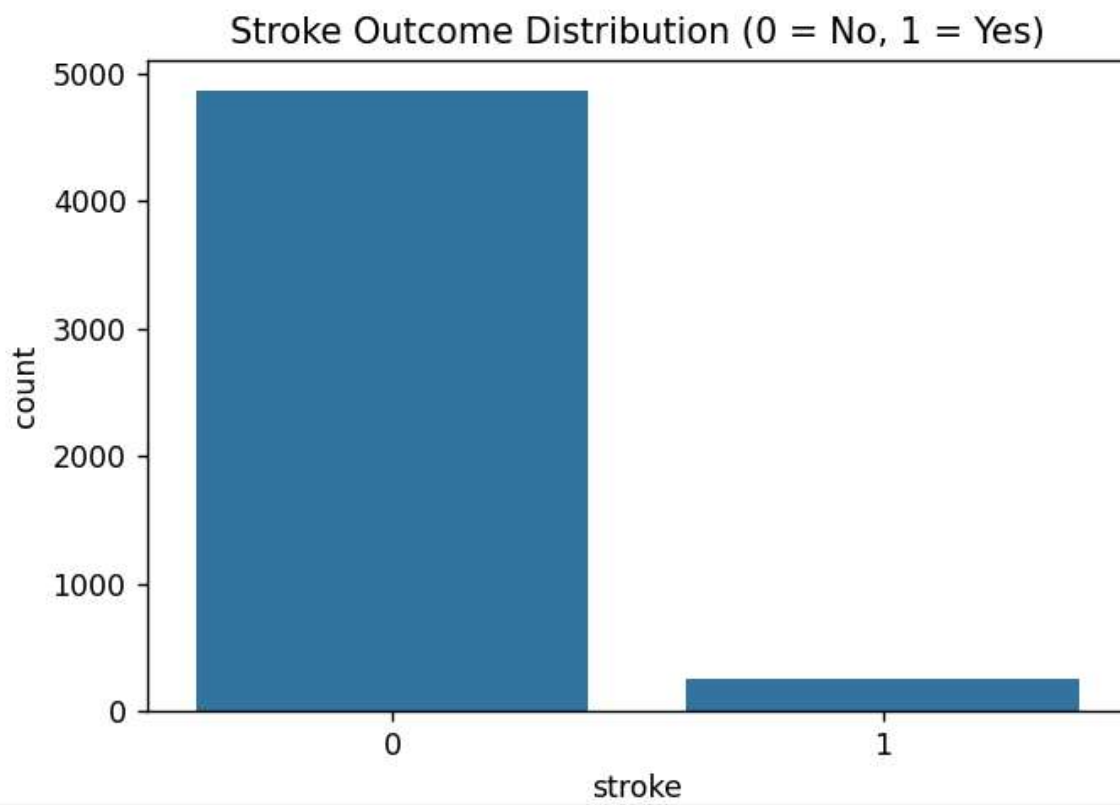
Missing values in each column:
id                0
gender            0
age              0
hypertension     0
heart_disease    0
ever_married     0
work_type        0
Residence_type   0
avg_glucose_level 0
bmi              201
smoking_status   0
stroke           0
dtype: int64

```


```

# 4. Target Variable Analysis
plt.figure(figsize=(6,4))
sns.countplot(x='stroke', data=df)
plt.title('Stroke Outcome Distribution (0 = No, 1 = Yes)')
plt.show()

```









```
# 5. Correlation Analysis (Numerical Features)
corr = df.corr(numeric_only=True)
plt.figure(figsize=(10,8))
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap')
plt.show()
```


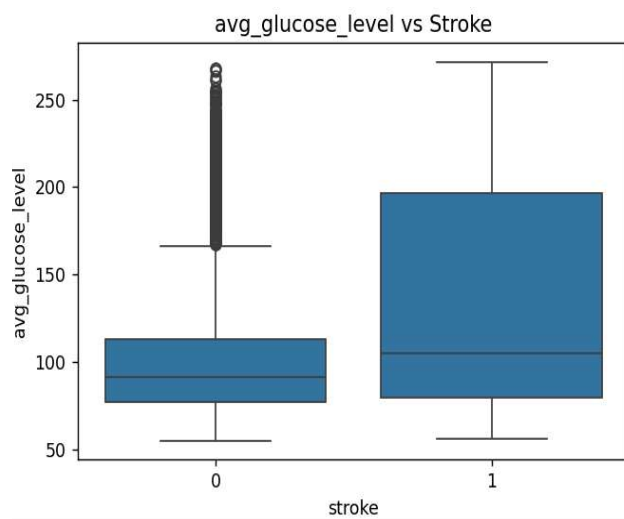

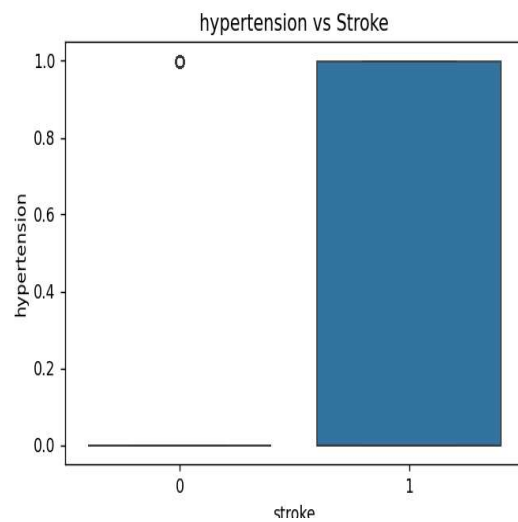



```
# 6. Feature Distribution Analysis
num_cols = df.select_dtypes(include=np.number).columns.tolist()
if 'id' in num_cols:
    num_cols.remove('id') # Remove 'id' as it's not a feature

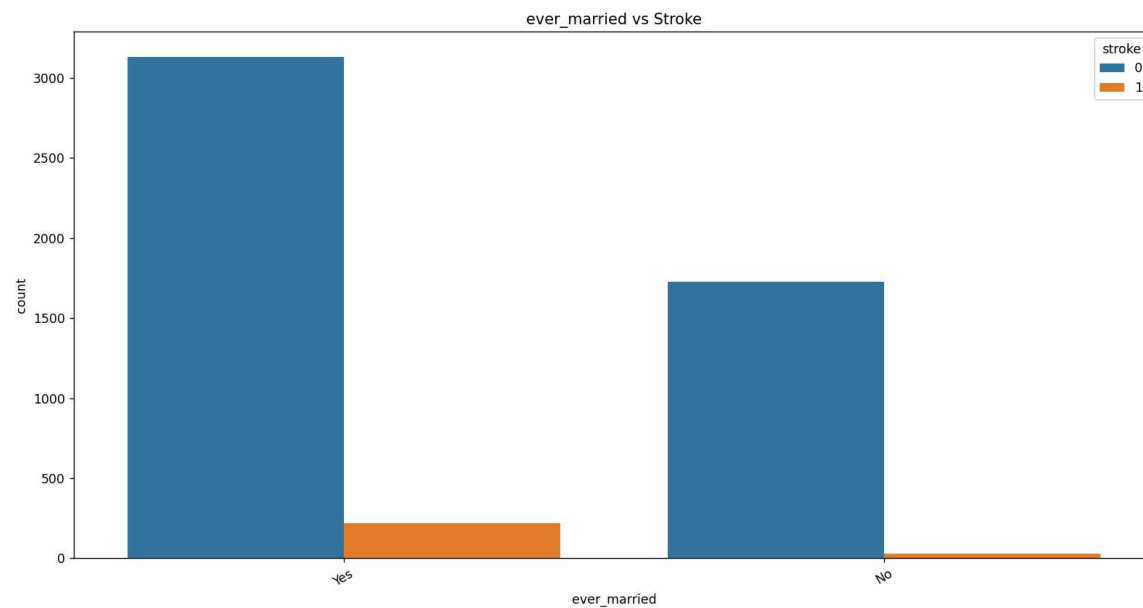

df[num_cols].hist(figsize=(12,8))
plt.suptitle("Histograms of Numerical Features", y=1.02)
plt.tight_layout()
plt.show()

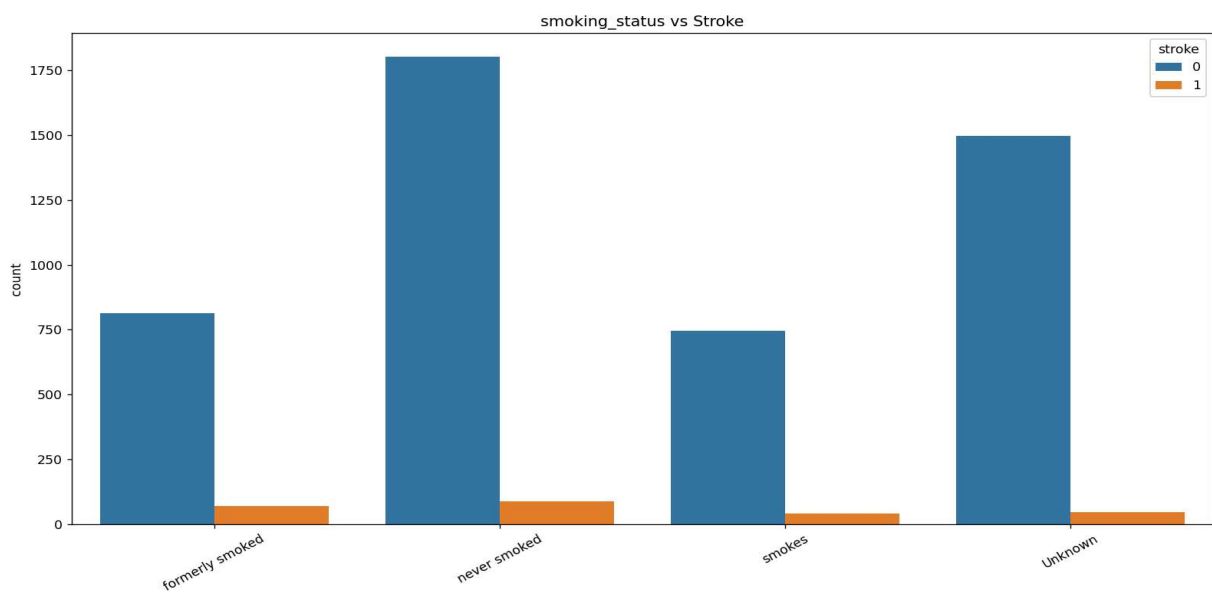
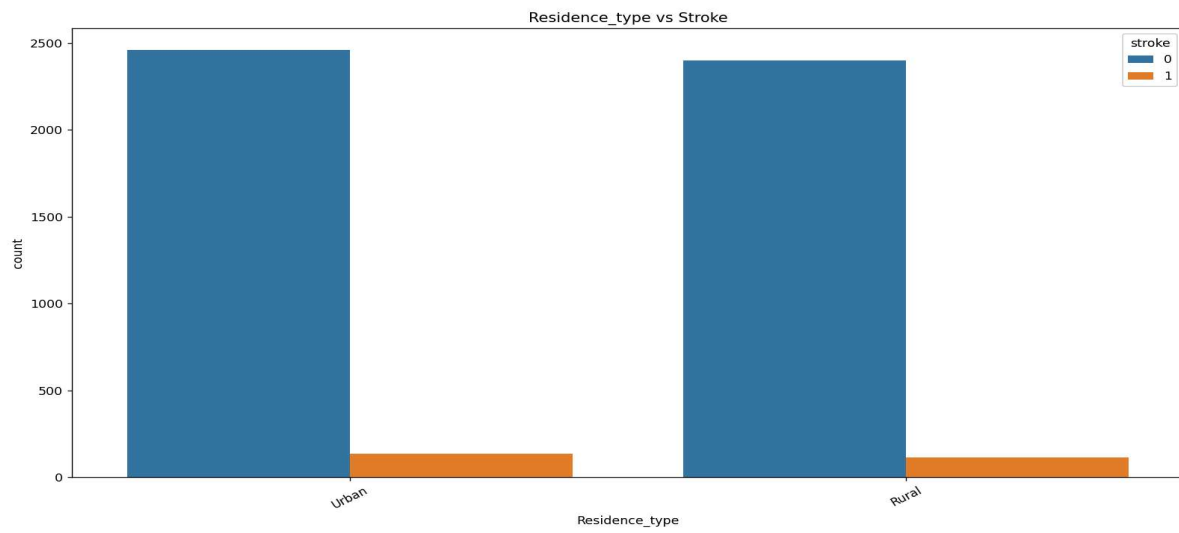
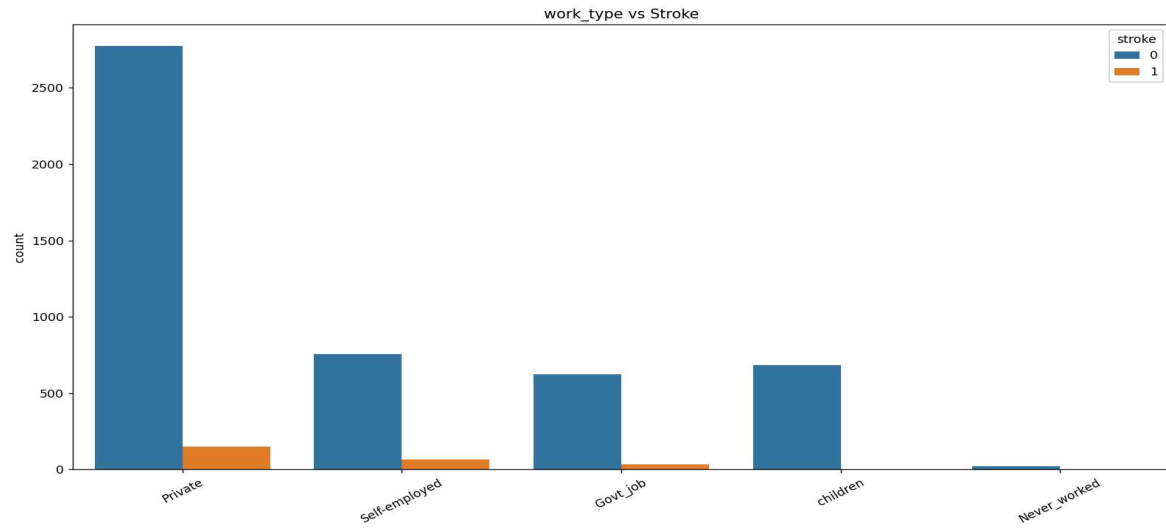
# Boxplots by Target
for col in num_cols:
    if col != 'stroke':
        plt.figure(figsize=(6,4))
        sns.boxplot(x='stroke', y=col, data=df)
        plt.title(f'{col} vs Stroke')
        plt.show()
```





```
# 7. Categorical Feature Analysis
cat_cols = [col for col in cat_cols if col != 'id'] # Exclude 'id' if present
for col in cat_cols:
    plt.figure(figsize=(7,4))
    sns.countplot(x=col, hue='stroke', data=df)
    plt.title(f'{col} vs Stroke')
    plt.xticks(rotation=30)
    plt.show()
```





```
# 8. Multivariate Visualization (Pairplot)
selected_features = ['age', 'avg_glucose_level', 'bmi', 'stroke']
sns.pairplot(df[selected_features].dropna(), hue='stroke', palette='Set2', diag_kind='kde')
plt.suptitle("Pairplot of Selected Features", y=1.02)
plt.show()

# 9. Example: Grouped Analysis (Age Groups)
df['age_group'] = pd.cut(df['age'], bins=[0,20,40,60,80,100,120], labels=['0-20', '21-40', '41-60', '61-80', '81-100', '100+'])
plt.figure(figsize=(8,4))
sns.countplot(x='age_group', hue='stroke', data=df)
plt.title('Age Group vs Stroke')
plt.show()

# 10. Summary Statement (Edit or expand as needed)
print("""
Interpretation & Summary:
- Review the class balance for the target variable 'stroke'.
- Identify which features differ between stroke and non-stroke groups.
- Check for missing values and outliers (especially in 'bmi').
- Assess correlation strength between features and the target.
- Use these insights for preprocessing and modeling.
""")
```

